

Multinomial Discrete Choice: Mixed Logit

Chris Conlon

Fall 2026

Grad IO

Mixed/ Random Coefficients Logit

An alternative is to allow for individuals to have **random coefficients** β_i , mixing over various logits:

$$\begin{aligned}\beta_i &= \bar{\beta} + \Pi y_i + \Sigma \nu_i & \mu_{ij} &= x_j \cdot (\Pi y_i + \Sigma \nu_i) \\ u_{ijt} &= x_j \beta_i + \varepsilon_{ij} & &= x_j \beta + \mu_{ij} + \varepsilon_{ij} \\ \sigma_j(\theta) &= \int \frac{\exp[x_j \beta_i]}{1 + \sum_k \exp[x_k \beta_i]} f(\beta_i | \theta) & &= \int \frac{\exp[x_j \beta + \mu_{ij}]}{1 + \sum_k \exp[x_k \beta + \mu_{ik}]} f(\mu_i | \theta)\end{aligned}$$

- ▶ x_j are observed per usual, y_i are individual demographics/features, and ε_{ij} is IID Type I EV. The key is that ν_i is unobserved and mean zero (often normally distributed).
- ▶ Each individual draws a vector β_i / μ_i (separately from ε_i). Conditional on β_i / μ_i each person follows an IIA logit model.
- ▶ However we integrate (or mix) over many such individuals giving us a **mixed logit** or **heirarchical model** (if you are a statistician)
- ▶ It helps to think about fixing β_i / μ_i first and then integrating out over ε_i

$$\sigma_{ij}(\theta) = \int \frac{\exp[x_j \beta_i]}{1 + \sum_k \exp[x_k \beta_i]} f(\beta_i | \theta) \approx \sum_{s=1}^S w_i^s \frac{\exp[x_j \beta_i^s]}{1 + \sum_k \exp[x_k \beta_i^s]}$$

- We can only **approximate** the integral with **weights** w_i^s and **nodes** β_i^s
- We can allow for there to be two types of β_i in the population (high-type, low-type). **latent class model**.
 - We can allow β_i to follow an independent normal distribution for each component of x_{ij} such as $\beta_i = \bar{\beta} + \Sigma \nu_i$.
 - We can allow for correlated normal draws $\beta_i \sim N(\mu_\beta, \Sigma_\beta)$.
 - Can allow for non-normal distributions too (lognormal, exponential).
 - Why is normal so easy?

- ▶ The structure is extremely flexible but at a cost.
- ▶ We generally must perform the integration numerically.
- ▶ High-dimensional numerical integration is difficult. In fact, integration in dimension 8 or higher makes me very nervous.
- ▶ We need to be parsimonious in how many variables have unobservable heterogeneity.
- ▶ Again observed heterogeneity does not make life difficult so the more of that the better!
 - If we see individual income, education, distance to hospital, etc. we can always interact that with observed characteristics without doing any additional integration.

How does it work?

- ▶ Well we are mixing over individuals who conditional on β_i or μ_i follow logit substitution patterns, however they may differ wildly in their σ_{ij} and hence their substitution patterns.
- ▶ For example if we are buying cameras: I may care a lot about price, you may care a lot about megapixels, and someone else may care mostly about zoom.
- ▶ The basic idea is that we need to explain the heteroskedasticity of $Cov(\epsilon_i, \epsilon_\ell)$ what random coefficients do is let us use a basis from our X 's.
- ▶ If our X 's are able to span the space effectively, then an RC logit model can approximate any arbitrary RUM (such as probit) (McFadden and Train 2002).
- ▶ Of course if you have 1000 products and two random coefficients, you are asking for a lot.

At the level of an individual substitution patterns follow a plain logit:

$$\epsilon_{\sigma_j, p_j} = \frac{\partial \sigma_j}{\partial p_j} \cdot \frac{p_j}{\sigma_j} = \frac{p_j}{\sigma_j} \cdot \int \frac{\partial \sigma_{ij}}{\partial p_j} dF_i = \frac{p_j}{\sigma_j} \int \beta_i \cdot \sigma_{ij} \cdot (1 - \sigma_{ij}) dF_i$$

$$\epsilon_{\sigma_j, p_k} = \frac{\partial \sigma_j}{\partial p_k} \cdot \frac{p_k}{\sigma_j} = \frac{p_k}{\sigma_j} \cdot \int \frac{\partial \sigma_{ij}}{\partial p_k} dF_i = -\frac{p_k}{\sigma_j} \int \beta_i \cdot \sigma_{ij} \cdot \sigma_{ik} dF_i$$

Notation: $\sigma_{ij}(\boldsymbol{\mu}_i) \equiv \sigma_{ij}$ is the individual choice probability and $\sigma_j = \int \sigma_{ij} dF_i$ is its aggregate (market-level) counterpart.

At the level of an individual substitution patterns follow a plain logit:

$$D_{jk,i}(x) = \frac{\frac{\partial \sigma_{ik}}{\partial p_j}}{\left| \frac{\partial \sigma_{ij}}{\partial p_j} \right|} = \frac{\sigma_{ik}(x)}{1 - \sigma_{ij}(x)}$$

But at the aggregate level, we mix over heterogeneous individuals:

$$D_{jk}(x) = \int D_{jk,i}(x) w_i(z_j, z'_j, x) dF_i \quad \text{where} \quad w_i(z_j, z'_j, x) = \frac{q_{ij}(z'_j, x) - q_{ij}(z_j, x)}{q_j(z'_j, x) - q_j(z_j, x)}$$

Different interventions (price changes, product removals, quality changes, etc.) give different weighting schemes.

How do we estimate these models? (Derived from multinomial log-likelihood):

$$\theta_{MLE} = \arg \min_{\theta} - \sum_{i=1}^N \sum_{j=1}^J y_{ij} \ln \hat{\sigma}_{ij}(\theta)$$

- ▶ We need to perform numerical integration to get $\hat{\sigma}_{ij}(\theta)$ and its derivative: $\widehat{\frac{\partial \sigma_{ij}}{\partial \theta}}$.
- ▶ Consistency
 - Technically for fixed number of MC draws the MSLE estimator is inconsistent. Why?
 - Very accurate integration is even more important!

Can also do a simulated GMM estimator with moment conditions:

$$\mathbb{E} \left[(y_{ij} - \hat{\sigma}_{ij}(\theta)) | z_{ij} \right] = 0$$

- ▶ We need to perform numerical integration to get $\hat{\sigma}_{ij}(\theta)$ and its derivative: $\widehat{\frac{\partial \sigma_{ij}}{\partial \theta}}$.
- ▶ The **optimal instruments** in the sense of Chamberlain (1987) or Amemiya (1977) are the derivative of moments with respect to parameters (**scores**): $z_{ij} = \frac{\partial \log \sigma_{ij}}{\partial \theta}(\theta)$. (Much more on this when we get to BLP instruments next week.)
- ▶ The true scores are infeasible (they depend on θ_0).
 - At true scores: same FOC as MLE.
 - Simulated scores have same consistency issue as MSLE.

As an alternative we could work with the score function directly:

$$\sum_{i=1}^n \frac{\partial \widehat{\log \sigma_{ij}}}{\partial \theta}(\theta) = \sum_{i=1}^N \frac{1}{\widehat{\sigma_{ij}}(\theta)} \cdot \frac{\partial \widehat{\sigma_{ij}}}{\partial \theta}(\theta) = 0$$

- ▶ We can get unbiased estimated of derivative (linearity)
- ▶ We can't get unbiased estimate of $\frac{1}{\widehat{\sigma_{ij}}(\theta)}$
 - Train's book suggests Accept-Reject
 - Maybe importance sampling would work?

How bad is the simulation error?

- ▶ Depends how small your shares are.
- ▶ Since you care about $\log \sigma_{jt}$ when shares are small, tiny errors can be enormous.
- ▶ Often it is pretty bad.

Estimation Details: Recommendations

You should read these separate notes, but my recommendations are:

1. Numerical integration:

- Use Gauss-Hermite quadrature rules to integrate over normal densities.
- Quadrature is great in low dimensions – but scales badly in high dimensions: if we need N_a points to accurately approximate the integral in $d = 1$ then we need N_a^d points in dimension d (using the tensor product of quadrature rules).
- Follow Heiss and Winschel (JoE) and do **sparse grids** in medium dimensions
- Try <http://sparse-grids.de>.

2. Nonlinear optimization: with mixtures the problem is non-convex (ie: very hard)

- You **must** provide analytic gradients.
- You should use a quasi-Newton solver and may need to try several.
- Try multiple starting values.

3. Generalized Method of Moments

- You can read more about optimal IV in these settings.